

# **Machine Learning for Link Prediction**

---

Giacomo Fiumara

# Table of contents

1. Introduction
2. Dimensionality Reduction: PCA and t-SNE
  - Principal Component Analysis (PCA)
  - t-SNE: t-Distributed Stochastic Neighbor Embedding
3. Random Walks on Networks
4. Word2Vec: Learning Word Embeddings from Text
  - CBOW: Continuous Bag of Words
  - Skip-Gram
  - Comparison and Extensions
5. Embedding Methods for Link Prediction
6. DeepWalk: Learning Social Representations

# Introduction

---

# Classical vs ML Approaches for Link Prediction

- **Classical heuristics** predict links using explicit formulas based only on graph topology.
- Examples include Common Neighbors, Jaccard Coefficient, Adamic-Adar, Preferential Attachment, Katz Index, and SimRank.
- These heuristics are transparent, easy to compute, and do not require training data.

# Classical vs ML Approaches for Link Prediction

## Limitations of classical heuristics

- Only use structural patterns; cannot combine with node or edge attributes.
- Can't learn complex patterns automatically; model is fixed regardless of data type.
- Performance is strongly dependent on network topology and may be poor on graphs that do not have high clustering or hub structure.

## Classical vs ML Approaches for Link Prediction


- **Machine Learning (ML) Approaches** leverage multiple features and learn optimal predictive combinations from data.
- ML models can incorporate classical heuristics as features, along with node/edge attributes and embedding-based signals.
- ML allows flexible modeling by training classifiers on labeled data (existing links and random non-existing links).

# Classical vs ML Approaches for Link Prediction

## Advantages of ML

- Learns from the actual data structure (adapts to domain specifics).
- Can easily incorporate node attributes, temporal information, and embeddings.
- Often achieves better performance than classical methods, especially for complex networks.

# Classical vs ML: Visual Comparison Example



Comparing CN predictions and ML classifier output on the same graph.

## Defining the ML Task for Link Prediction

- **Goal:** Given a network, predict which pairs of nodes may form (or should form) a link.
- ML approaches formulate link prediction as a supervised classification task.
- Each possible pair of nodes  $(x, y)$  is represented by a **feature vector** summarizing available information.

# Defining the ML Task for Link Prediction

## Common features

- Topological heuristics: Common Neighbors, Adamic-Adar, Jaccard, Preferential Attachment, Katz, etc.
- Node attributes: profile similarity, group membership, categorical features.
- Edge attributes: weights, timestamps, types (if available).
- Node embeddings: vector representations from algorithms like node2vec or DeepWalk.

## Defining the ML Task for Link Prediction

- Feature vector for a pair  $(x, y)$ :


$$X_{(x,y)} = [s_{CN}(x, y), s_{AA}(x, y), s_J(x, y), \text{attributes}, \text{embeddings}, \dots]$$

- Ground truth label for  $(x, y)$ :

$$Y_{(x,y)} = \begin{cases} 1 & \text{if edge exists (or is hidden for testing)} \\ 0 & \text{otherwise} \end{cases}$$

- The ML model is trained to classify feature vectors  $X_{(x,y)}$  as linked ( $Y = 1$ ) or not linked ( $Y = 0$ ).

# ML Link Prediction Task: Feature Representation



For every candidate node pair, construct a feature vector from topological and attribute-based measures.

## Train/Test Splitting in Link Prediction

- **Train/test splitting** is necessary to evaluate the ML model's ability to predict missing or future links.
- A common approach: randomly hide a fraction of true edges; these "hidden" links become the positive test set.
- Negative test samples are pairs of nodes not connected in the current network.

# Train/Test Splitting in Link Prediction

## ML workflow

- **Training set:** All not-hidden edges (positive), and an equal or larger number of random non-edges (negative).
- **Test set:** Hidden edges (positive), plus sampled non-edges (negative).

## Other splitting strategies

- **Temporal splitting:** Use earlier time steps for training, later for testing (important for evolving networks).
- **Node-induced splitting:** Hold out links associated with a subset of nodes to simulate cold-start prediction.

## Train/Test Splitting in Link Prediction


- Example: Let  $E_{train}$  be the set of observed edges,  $E_{test}$  be the set of hidden edges,  $N_{train}/N_{test}$  non-edges:

$$X_{train}, Y_{train} \quad (E_{train} = 1, N_{train} = 0)$$

$$X_{test}, Y_{test} \quad (E_{test} = 1, N_{test} = 0)$$

- It is crucial to ensure training and test sets do not overlap and reflect the prediction scenario of interest.

# Train/Test Splitting: Visual Example



Original network with held-out test links and candidate non-edges.

# **Dimensionality Reduction: PCA and t-SNE**

---

## Why Dimensionality Reduction?

- High-dimensional data (e.g., node embeddings, gene expression, images) are hard to visualize or interpret directly.
- Dimensionality reduction projects data to a lower-dimensional space (usually 2D or 3D) while preserving key patterns.
- Visualization tools like PCA and t-SNE help:
  - Find clusters, outliers, and structure.
  - Interpret and communicate complex models.

## Principal Component Analysis (PCA): Intuition

- **PCA** is a linear technique for dimensionality reduction and data visualization.
- Identifies directions (principal components) in data with the most variance.
- Projects data onto the first  $k$  components to obtain a summary with minimal information loss.
- Useful for projecting high-dimensional embeddings to 2D/3D for visualization.

## PCA: Mathematical Description

- Given data matrix  $X$  (rows are samples, columns are features), center and compute the covariance:

$$S = \frac{1}{n-1} X^T X$$


- Find eigenvectors  $\mathbf{w}_1, \mathbf{w}_2, \dots$  of  $S$ : directions with highest variance.
- Project data:

$$Z = X \cdot W_k$$

where  $W_k$  contains the top  $k$  eigenvectors as columns.

- Result: each sample is now a  $k$ -dimensional vector ( $k \ll D$ ).

# PCA: Example on a Toy Dataset




Data projected onto main axes of variance in 2D.

## PCA for Embedding Visualization

- PCA can be used to project node embeddings (e.g., from node2vec/DeepWalk) to 2D for direct visualization.
- Global structure (relative distances, spread) is largely preserved.
- Good for revealing clusters and overall geometric relationships in high dimensions.


# PCA for Embedding Visualization



## Why Do We Need t-SNE? Limitations of Linear Methods

- Linear methods like PCA are interpretable and preserve global structure, but they may fail to reveal complex, nonlinear patterns.
- In many real datasets—including network embeddings—clusters and groups are arranged in nonlinear manifolds that PCA cannot separate.
- Example: the "two moons" or "Swiss roll" datasets have clear geometric structure that PCA projects poorly, but nonlinear methods reveal easily.
- **t-SNE** is designed to discover and visualize such nonlinear structure by preserving local neighborhoods rather than global distances.

## Example: PCA Fails on Nonlinear Data



## t-SNE: The Core Idea

- t-SNE stands for **t-Distributed Stochastic Neighbor Embedding**.
- The algorithm models similarity between data points in high-dimensional space and then finds a low-dimensional (2D/3D) embedding that preserves these similarities.
- Key principle: **Points that are neighbors in high dimensions should remain neighbors in low dimensions.**
- Unlike PCA, t-SNE is **nonlinear** and does not preserve global distances—only local structure matters.

## t-SNE Step 1: High-Dimensional Similarities

- In the original high-dimensional space, compute pairwise similarities between points using a **Gaussian kernel**:

$$p_{j|i} = \frac{\exp\left(-\frac{\|x_i - x_j\|^2}{2\sigma_i^2}\right)}{\sum_{k \neq i} \exp\left(-\frac{\|x_i - x_k\|^2}{2\sigma_i^2}\right)}$$

- Here,  $p_{j|i}$  is the probability that point  $i$  "picks" point  $j$  as a neighbor in high dimensions.
- $\sigma_i$  is a bandwidth parameter (often set via a target "perplexity"—roughly the expected number of neighbors).
- Smaller distances yield higher  $p_{j|i}$ ; distant points have near-zero probability.

## Perplexity: Controlling Neighborhood Size

- **Perplexity** is a hyperparameter (e.g., 5, 30, 50) that controls the expected size of each point's neighborhood.
- For each point  $i$ , we set  $\sigma_i$  such that the entropy of the conditional distribution  $p_{j|i}$  equals  $\log(\text{perplexity})$ .
- Typical range: 5–50. Higher perplexity considers larger neighborhoods (more global structure); lower perplexity focuses on very local structure.
- Choose perplexity based on dataset size: a rule of thumb is  $\text{perplexity} \approx \sqrt{n}$ , where  $n$  is the number of samples.

## t-SNE Step 2: Low-Dimensional Similarities

- In the low-dimensional space (2D or 3D), place points  $y_1, \dots, y_n$  and define similarities using a **heavy-tailed Student-t distribution**:

$$q_{j|i} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq i} (1 + \|y_i - y_k\|^2)^{-1}}$$

- Why Student-t? It has heavier tails than a Gaussian, which helps avoid overcrowding of points in 2D (the "crowding problem").
- The form  $(1 + \|y_i - y_j\|^2)^{-1}$  means even moderately distant points in 2D have non-negligible similarity.

## t-SNE Step 3: Minimize KL Divergence

- t-SNE finds  $y_1, \dots, y_n$  that minimize the Kullback-Leibler (KL) divergence between high-dimensional and low-dimensional similarity distributions:

$$\text{KL}(P\|Q) = \sum_i \sum_j p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

- Here,

$$p_{ij} = \frac{1}{2}(p_{j|i} + p_{i|j})$$

and

$$q_{ij} = \frac{1}{2}(q_{j|i} + q_{i|j})$$

are symmetrized similarities.

## t-SNE Step 3: Minimize KL Divergence

- Minimization is done via gradient descent (or adaptive methods like Adam).
- Small KL means the 2D layout preserves neighborhood structure: if  $i$  and  $j$  are close in high-D, they remain close in 2D.

## Gradient of KL Divergence

- The gradient of KL with respect to  $y_i$  is:


$$\frac{\partial \text{KL}}{\partial y_i} = 4 \sum_j (p_{ij} - q_{ij})(y_i - y_j)(1 + \|y_i - y_j\|^2)^{-1}$$

- Intuition: if  $p_{ij} > q_{ij}$  (points are neighbors in high-D but far in low-D), the gradient pulls  $y_i$  and  $y_j$  together.
- If  $p_{ij} < q_{ij}$  (unrelated in high-D, close in low-D), the gradient pushes them apart.
- Optimization alternates gradient updates with "momentum" and learning rate annealing.


# t-SNE: The Complete Algorithm

1. **Input:** High-dimensional data  $X \in \mathbb{R}^{n \times d}$ , perplexity  $\text{perp}$ , number of iterations  $T$ , learning rate  $\eta$ .
2. Compute high-D pairwise similarities  $p_{ij}$  (Gaussian kernel, adjust  $\sigma_i$  for each point to match perplexity).
3. Initialize  $y_i \in \mathbb{R}^2$  randomly (e.g., from  $\mathcal{N}(0, 10^{-4}I)$ ).
4. For iteration  $t = 1, \dots, T$ :
  - 4.1 Compute low-D similarities  $q_{ij}$  (Student-t).
  - 4.2 Compute KL divergence and its gradient.
  - 4.3 Update  $y_i$  via gradient descent with momentum.
5. **Output:** Low-dimensional embeddings  $Y = \{y_1, \dots, y_n\}$ .

## Example: t-SNE on Two Moons



# t-SNE on a Network Embedding: Karate Club




## Advantages of t-SNE

- **Reveals nonlinear structure:** Clusters and manifolds are visualized clearly.
- **Local neighborhood preservation:** Points that are neighbors in high-D remain neighbors in 2D.
- **Handles complex geometries:** Swiss roll, intertwined spirals, and other complex manifolds are straightened and separated.
- **Intuitive interpretation:** Close points in the plot are similar; distant points are different (locally).


## Limitations

- **Global structure is not preserved:** Distances between far-apart clusters are not meaningful.
- **Computationally expensive:**  $O(n^2)$  time and memory for  $n$  points; approximations for large datasets.
- **Hyperparameter sensitive:** Perplexity choice affects results; no single "right" perplexity for all datasets.
- **Non-deterministic:** Different random initializations and optimization runs can yield different (but topologically similar) layouts.


## Perplexity Comparison 1/3



## Perplexity Comparison 2/3



## Perplexity Comparison 3/3




## Practical Recommendations

- Use t-SNE when interested in discovering clusters and local neighborhoods.
- Set perplexity between 5 and 50; for most datasets, 30 is a good starting point.
- Always visualize with multiple perplexities to ensure stability.
- Do NOT interpret global distances or cluster separation sizes as meaningful—focus on local groupings.
- Combine t-SNE with other visualizations (PCA, UMAP, original network layout) for a complete picture.


## PCA vs t-SNE: Visualization Comparison

- **PCA preserves global structure:** relative distances in 2D reflect high-dimensional distances.
- **t-SNE preserves local neighborhoods:** clusters are faithfully preserved, but global layout can be stretched or compressed.

# PCA vs t-SNE: Visualization Comparison



# PCA vs t-SNE: Visualization Comparison



# **Random Walks on Networks**

---

## What is a Random Walk?

- A **random walk** is a stochastic process that starts at a node and, at each step, moves to a random neighbor.
- Formally, given a current node  $v$  with neighbors


$$\Gamma(v) = \{u_1, u_2, \dots, u_k\}$$

- The next node is chosen uniformly at random:

$$P(\text{next} = u_i \mid \text{current} = v) = \frac{1}{|\Gamma(v)|}$$

- A walk of length  $L$  is a sequence of nodes:  $(v_0, v_1, v_2, \dots, v_L)$  where each transition is governed by this random process.

## Simple Example: Random Walk on a Toy Graph



## Why Random Walks Explore Network Structure

- Random walks naturally spend more time in densely-connected regions (communities) because there are more paths to traverse within dense areas.
- Conversely, sparsely-connected regions are visited less frequently because there are fewer paths.
- If two nodes are in the same community (densely connected), they are more likely to co-occur in random walks.
- Random walks thus provide a way to **sample neighborhoods** that reflects network structure without explicitly computing distances or metrics.
- This is important for methods like DeepWalk and node2vec.

# Mathematical Properties of Random Walks

- **Stationary distribution:** For a connected, undirected graph, a random walk converges to a stationary distribution  $\pi(v) \propto \deg(v)$  (degree distribution).
- This means: in the long run, the probability of being at node  $v$  is proportional to its degree. High-degree nodes are visited more often.
- **Mixing time:** The number of steps needed to approach the stationary distribution. Fast mixing means the walk explores the graph quickly.
- **Commute time:** The expected time to go from node  $u$  to node  $v$  and back. Short commute times indicate  $u$  and  $v$  are in the same community.
- These properties explain why random walk based embeddings

# Generating Random Walks: Algorithm

```
import networkx as nx
import numpy as np

def random_walk(G, start_node, walk_length, seed=None):
    if seed is not None:
        rng = np.random.default_rng(seed)
    else:
        rng = np.random.default_rng()

    walk = [start_node]
    current = start_node

    for x in range(walk_length):
        neighbors = list(G.neighbors(current))
        if not neighbors:
            break
        current = rng.choice(neighbors)
        walk.append(current)

    return walk
```

# Generating Random Walks: Algorithm

```
#G = nx.cycle_graph(6) # nodes 0..5 in a ring
G = nx.karate_club_graph()

w = random_walk(G, start_node=0, walk_length=10, seed=42)
print("Random walk:", w)

def generate_walks(G, num_walks_per_node=5, walk_length=10, seed=None):
    walks = []
    rng = np.random.default_rng(seed)
    for node in G.nodes():
        for _ in range(num_walks_per_node):
            s = int(rng.integers(0, 1e9))
            walks.append(random_walk(G, start_node=node, walk_length=walk_length, seed=s))
    return walks


walks = generate_walks(G, num_walks_per_node=3, walk_length=8, seed=123)
for i, w in enumerate(walks[:5]):
    print(f"Walk {i}:", w)
```

```
Random walk: [0, 2, 27, 24, 27, 23, 33, 9, 33, 14, 32]
Walk 0: [0, 17, 0, 21, 1, 13, 3, 0, 12]
Walk 1: [0, 3, 1, 7, 2, 0, 17, 1, 17]
Walk 2: [0, 10, 5, 10, 4, 0, 6, 16, 5]
Walk 3: [1, 3, 7, 0, 8, 33, 26, 33, 31]
Walk 4: [1, 7, 2, 0, 11, 0, 19, 0, 1]
```

## Multiple Random Walks: Sampling Neighborhoods at Scale

- One random walk from a node provides a single sample of its neighborhood.
- Many walks from each node sample different neighborhoods and capture differences in network structure.
- Each walk explores a different path, discovering neighbors at different distances and in different communities.
- Longer walks explore farther, capturing global structure.  
Shorter walks stay more local, focusing on immediate neighborhoods.
- A collection of walks of various lengths captures a rich, multi-scale view of the network neighborhood.

# Random Walks Visualized: Distribution of Visited Nodes



## From Walks to Embeddings

- If node  $u$  and node  $v$  often appear together in random walks, they should have similar embeddings.
- This is the foundation of embedding methods:
  - Generate walks to collect neighborhoods.
  - Apply a language model (skip-gram) to learn embeddings from walk sequences.
  - The language model sees walks as sentences and learns to represent co-occurring nodes as similar vectors.
- This two-step approach (walks + skip-gram) is exactly what DeepWalk and node2vec do.

# **Word2Vec: Learning Word Embeddings from Text**

---

## What is Word2Vec?

- **Word2Vec** learns a real-valued vector embedding for each word in a vocabulary.
- Key insight: *Words that appear in similar contexts should have similar vectors.*
- Training signal comes from the text itself:
  - Look at a word and its neighbors in a sliding window.
  - Define a prediction task ( $\text{context} \rightarrow \text{word}$ , or  $\text{word} \rightarrow \text{context}$ ).
  - Learn embeddings that are good at this task.
- Result: word vectors capture semantic and syntactic relationships.

## The Two Flavors: CBOW vs Skip-Gram

- CBOW (Continuous bag of words) and Skip-Gram are **shallow neural networks** with the same architecture.
- They differ in the **direction of prediction**:

	CBOW	Skip-Gram
Input	Context words	Target word
Output	Target word	Context words
Direction	context → target	target → context

- Despite the swap, the learned embeddings capture the same kind of similarity.


## CBOW: High-Level Idea

- **Goal:** Predict a target word from its surrounding context words.
- **Example:** In the sentence “the quick brown fox jumps over”, with window size 2:
  - Context: {quick, brown, jumps, over}
  - Target: fox

# CBOW: High-Level Idea

- Why “continuous”?
  - Use dense real-valued vectors (embeddings) in a continuous space.
  - Not sparse one-hot symbols.
- Why “bag”?
  - Order of context words does not matter; they are summed/averaged.
  - Only the multiset (bag) of words in the window matters.

## CBOW: context



## CBOW: Mathematical Formulation

- Let  $\mathbf{v}_w \in \mathbb{R}^d$  be the embedding of word  $w$ .
- Context words:  $\{w_{-m}, w_{-m+1}, \dots, w_{-1}, w_1, \dots, w_m\}$  (window of size  $m$  on each side).
- **Step 1:** Average context embeddings:

$$\mathbf{h} = \frac{1}{2m} \sum_{i \in \text{context}} \mathbf{v}_i$$

## CBOW: Mathematical Formulation

- **Step 2:** Predict target word  $w_t$  via softmax:

$$p(w_t \mid \text{context}) = \frac{\exp(\mathbf{v}'_{w_t} \cdot \mathbf{h})}{\sum_{w \in V} \exp(\mathbf{v}'_w \cdot \mathbf{h})}$$

where  $\mathbf{v}'_w$  is the “output embedding” of word  $w$ .

- **Loss:** Minimize cross-entropy (equivalently, maximize log-likelihood of true target word).

# CBOW: Strengths and Weaknesses

## Strengths:

- *Fast to train.* Fewer training examples needed per update (averages many contexts).
- Works well for *frequent* words.
- Simpler architecture.

## Weaknesses:

- Less effective for *rare* words (each rare word is center less often).
- Loses information by averaging (order of context is ignored).
- Often produces slightly lower-quality embeddings than skip-gram.

## Skip-Gram: High-Level Idea

- **Goal:** Predict context words from a target word.
- **Example:** In the sentence “the quick brown fox jumps over”, with window size 2:
  - Target: fox
  - Create pairs:  
 $(fox, quick)$ ,  $(fox, brown)$ ,  $(fox, jumps)$ ,  $(fox, over)$
- **Training:** For each pair, train to predict the context word from the target.
- **Intuition:** Learn embeddings where the target word’s vector is good at “pointing towards” its context neighbors.

## Skip-Gram: Mathematical Formulation

- Let  $\mathbf{v}_w$  be the embedding of word  $w$ , and  $\mathbf{v}'_w$  be its output embedding.
- Target word:  $w_t$ , context word:  $w_c$  (a neighbor within the window).
- Step 1:** Embed target word:

$$\mathbf{h} = \mathbf{v}_{w_t}$$

- Step 2:** For each context word  $w_c$ , predict it via softmax:

$$p(w_c | w_t) = \frac{\exp(\mathbf{v}'_{w_c} \cdot \mathbf{h})}{\sum_{w \in V} \exp(\mathbf{v}'_w \cdot \mathbf{h})}$$

# Skip-Gram: Strengths and Weaknesses

## Strengths:

- Works very well for *rare* words (each rare word used as target many times).
- Often produces *higher-quality* embeddings.
- Better at capturing fine-grained relationships.

## Weaknesses:

- *Slower to train.* More training pairs per sentence.
- Requires careful regularization.
- Negative sampling trick is needed for practical speed.

## CBOW vs Skip-Gram at a Glance

Aspect	CBOW	Skip-Gram
Direction	Context → Word	Word → Context
Training speed	Fast	Slow
Best for	Frequent words	Rare words
Quality	Good	Often better
Typical use	Speed-critical	Quality-critical

## Negative Sampling Trick

- Full softmax over vocabulary is expensive (compute probabilities for all  $|V| \approx 10^6$  words).
- **Negative sampling:** Replace softmax with binary classification.
- For each training pair (target, context):
  - Positive example: (target, context word)
  - Negative examples: sample  $k$  random words from vocabulary as “noise”

## Negative Sampling Trick

- Train a binary classifier: “Is this word a true context neighbor, or noise?”
- Much faster:  $O(k)$  instead of  $O(|V|)$  per update, with  $k \ll |V|$  (e.g.,  $k = 5$ ).
- Empirically: produces very similar embeddings to full softmax.

## Connection to Node2Vec and DeepWalk

- Word2Vec (skip-gram) learns embeddings for words in text, where:
  - “Sentences” = text sequences.
  - “Words” = vocabulary items.
  - “Context” = neighboring words in a sliding window.
- **DeepWalk/Node2Vec:** Apply the exact same skip-gram idea to graphs:
  - “Sentences” = random walks on the graph.
  - “Words” = nodes (or graph entities).
  - “Context” = neighboring nodes along the random walk.
- Result: nodes with similar network neighborhoods get similar vectors.
- **Key insight:** The embedding technique is domain-agnostic. It works if it is possible to define a notion of “context”.

## Key Takeaway

- CBOW and skip-gram are **two sides of the same coin**:
  - CBOW: context → word (faster, good for frequent words)
  - Skip-gram: word → context (slower, better for rare words)
- Both learn embeddings in a **continuous vector space** using **shallow neural networks**.
- The key intuition: *words (or nodes) appearing in similar contexts should have similar embeddings.*
- The method is **general**: it applies to any sequential or contextual data, including graphs (node2vec, DeepWalk).

## Word Analogies: Vector Arithmetic

- Word2Vec represents each word as a vector in a continuous space.
- An **analogy** is solved using **vector arithmetic**.
- **Example:** “king is to man as queen is to woman”
  - Compute:  $\mathbf{v}_{\text{king}} - \mathbf{v}_{\text{man}} + \mathbf{v}_{\text{woman}}$
  - Find the word whose vector is **closest** to this result (by cosine similarity).
  - Ideally: queen

## Word Analogies: Vector Arithmetic

- In code (Gensim), this is:
  - **Positive words:** vectors that are **added**. Example: `["king", "woman"]`
  - **Negative words:** vectors that are **subtracted**. Example: `["man"]`
- **Intuition:**
  - Positive = capture properties you want in the result (“king-like” and “woman-like”).
  - Negative = remove unwanted properties (subtract “man-ness”).

## Why Vector Arithmetic Works: Learned Semantic Directions

- Word2Vec is not explicitly told to create a “gender axis” or “royalty axis”.
- However, during training, **systematic relationships** emerge naturally:
  - The difference vector  $\mathbf{v}_{\text{woman}} - \mathbf{v}_{\text{man}}$  encodes a “gender shift”.
  - This same direction applies across different word pairs.
- Key observation:**

$$\mathbf{v}_{\text{woman}} - \mathbf{v}_{\text{man}} \approx \mathbf{v}_{\text{queen}} - \mathbf{v}_{\text{king}}$$

- Therefore:

$$\mathbf{v}_{\text{king}} - \mathbf{v}_{\text{man}} + \mathbf{v}_{\text{woman}} \approx \mathbf{v}_{\text{queen}}$$

# **Embedding Methods for Link Prediction**

---

## Embedding Methods: Motivation and Context

- Embedding methods transform nodes into vectors in a continuous space, capturing structural similarity and patterns in the network.
- The goal is to map each node  $i$  to a vector  $\mathbf{z}_i \in \mathbb{R}^d$  such that nodes with similar network positions are close in the embedding space.
- Embeddings can capture higher-order relationships and community structure that are missed by purely local heuristics.
- Once nodes are embedded, link prediction is reduced to estimating the likelihood of links via vector operations (e.g., dot product, cosine similarity).

## Node2Vec and DeepWalk: Random Walk-Based Embeddings

- **Node2Vec** and **DeepWalk** are popular algorithms that use short random walks to generate "sentences" over the network, then apply language modeling techniques to learn embeddings.
- Random walk-based embeddings (e.g., DeepWalk) treat all neighbors equally: at each step, pick a random neighbor uniformly.
- This captures both local and global structure, but without explicit control over what type of structure is emphasized.

## Motivation: Beyond Unbiased Random Walks

- Different applications may need different biases:
  - Homophily-based prediction: nodes in the same community should be similar.
  - Structural equivalence: nodes with similar roles (e.g., bridges, hubs) should be similar, even if far apart.
- **node2vec** introduces **biased random walks** via two parameters,  $p$  and  $q$ , to interpolate between these goals.

## node2vec Overview

- How **node2vec** learns vector embeddings for nodes:
  1. Generate biased random walks starting from each node.
  2. Treat walks as "sentences" and train a skip-gram language model (word2vec-style).
  3. Output: a dense vector for each node.
- Nodes that frequently co-occur in walks (via the biased sampling) end up close in embedding space.
- The bias parameters  $p$  and  $q$  allow us to control whether embeddings emphasize local communities or global structural roles.

# The Optimization Objective

- For each node  $u$ , node2vec learns embeddings to maximize the log-probability of observing its walk-based neighborhood  $N_S(u)$ :

$$\max_{\{\mathbf{z}_u\}} \sum_{u \in V} \log P(N_S(u) | \mathbf{z}_u)$$

- Expanding the conditional probability using the chain rule:

$$P(N_S(u) | \mathbf{z}_u) = \prod_{v \in N_S(u)} P(v | \mathbf{z}_u)$$

- Using a softmax (or negative sampling in practice):

$$P(v | \mathbf{z}_u) = \frac{\exp(\mathbf{z}_v^\top \mathbf{z}_u)}{\sum_{w \in V} \exp(\mathbf{z}_w^\top \mathbf{z}_u)}$$

## What is $N_S(u)$ ? The Walk-Based Neighborhood

- $N_S(u)$  is the **multiset of nodes that co-occur with  $u$  in random walk contexts**, sampled according to strategy  $S$ .
- Concretely:
  - From node  $u$ , generate multiple random walks:  $(u = v_0, v_1, \dots, v_L)$ .
  - Use a context window of size  $k$  (e.g.,  $k = 5$ ).
  - For node  $v_i$  at position  $i$ , its context nodes are  $v_{i-k}, \dots, v_{i-1}, v_{i+1}, \dots, v_{i+k}$  (within walk bounds).
- $N_S(u)$  aggregates all context nodes from all walks starting at  $u$ .

## Walk-Based Neighborhood: Example on a Toy Graph

- Consider a simple path graph:  $0 - 1 - 2 - 3 - 4$ .
- Starting from node 2, suppose we generate a walk  $(2, 1, 0, 1, 2, 3, 4)$ .
- With context window  $k = 2$ :
  - At position 0 (node 2): context =  $\{1, 0\}$ .
  - At position 4 (node 2): context =  $\{0, 1, 3, 4\}$ .
  - $N_S(2)$  collects these and all contexts from other walks started at 2.

## Walk-Based Neighborhood: Example on a Toy Graph

- Notice:  $N_S(2)$  includes nodes beyond immediate neighbors (e.g., node 0, node 4), reflecting the walk-induced structure.
- This is analogous to word2vec: the "context" of a word reflects which other words appear nearby in sentences, not just dictionary neighbors.

## Analogy to Word2Vec

- In word2vec (language model):
  - "word" = vocabulary item (e.g., "cat", "dog").
  - "context" = words that frequently appear within a fixed window (e.g., 5 words left/right).
  - Learning objective: maximize  $\log P(\text{context} \mid \text{word})$  across training sentences.
- In node2vec:
  - "word" = node ID (e.g., node 0, node 1).
  - "sentence" = random walk (sequence of nodes).
  - "context" = nodes within window  $k$  of the current node in the walk ( $N_S(u)$ ).
  - Learning objective: maximize  $\log P(N_S(u) \mid \mathbf{z}_u)$  across all nodes and sampled walks.

## Second-Order Biased Random Walks: The Key Innovation

- In a plain random walk (DeepWalk), the transition from current node  $v$  to next node  $x$  depends only on  $v$ :

$$P(x \mid v) = \frac{1}{\deg(v)} \quad (\text{uniform})$$

- In node2vec, the transition depends on **both** the current node  $v$  **and** the previous node  $t$ :

$$P(x \mid t, v) \propto \alpha_{pq}(t, x) \cdot w_{vx}$$

## Second-Order Biased Random Walks: The Key Innovation

- This is called **second-order** because the transition depends on the last two nodes (previous + current).
- The bias factor  $\alpha_{pq}(t, x)$  encodes preferences based on the relationship between previous node  $t$  and candidate next node  $x$ .

## The Bias Factor $\alpha_{pq}(t, x)$ : Definition

- Given that a walk has just moved from  $t$  to current node  $v$ , and we are deciding whether to step to neighbor  $x$ :
- Let  $d_{tx}$  = shortest-path distance between  $t$  and  $x$  in the graph.
- Define the unnormalized bias weight:

$$\alpha_{pq}(t, x) = \begin{cases} \frac{1}{p} & \text{if } d_{tx} = 0 \quad (\text{return to } t) \\ 1 & \text{if } d_{tx} = 1 \quad (\text{stay near } t) \\ \frac{1}{q} & \text{if } d_{tx} = 2 \quad (\text{move away from } t) \end{cases}$$

## The Bias Factor $\alpha_{pq}(t, x)$ : Definition

- The actual transition probability is then normalized:

$$P(x \mid t, v) = \frac{\alpha_{pq}(t, x) \cdot w_{vx}}{\sum_{z \in \Gamma(v)} \alpha_{pq}(t, z) \cdot w_{vz}}$$

## Parameter $p$ : Return Parameter

- $p$  controls the probability of immediately returning to the previous node  $t$  (i.e., backtracking).
- If  $x = t$ , then  $d_{tx} = 0$  and  $\alpha_{pq}(t, x) = \frac{1}{p}$ .
- **Large  $p$**  (e.g.,  $p = 2$ ):
  - $\frac{1}{p}$  is small, so backtracking is **discouraged**.
  - The walk is more likely to keep moving forward into new regions.

## Parameter $p$ : Return Parameter

- **Small  $p$**  (e.g.,  $p = 0.5$ ):
  - $\frac{1}{p}$  is large, so backtracking is **encouraged**.
  - The walk spends more time exploring the local neighborhood around the current node.
  - Default:  $p = 1$  (no preference for or against backtracking).

## Parameter $q$ : In-Out Parameter

- $q$  controls the trade-off between local exploration (BFS-like) and global exploration (DFS-like).
- Nodes at distance  $d_{tx} = 2$  from  $t$  (i.e., moving away from  $t$ ) have weight  $\frac{1}{q}$ .
- **Large  $q$**  (e.g.,  $q = 2$ ):
  - $\frac{1}{q}$  is small, so moves away from  $t$  are **down-weighted**.
  - The walk stays clustered near the current region (BFS-like).
  - Good for discovering **homophilic communities**: nodes in the same community are visited together.

## Parameter $q$ : In-Out Parameter

- **Small  $q$**  (e.g.,  $q = 0.5$ ):
  - $\frac{1}{q}$  is large, so moves away from  $t$  are **up-weighted**.
  - The walk is more exploratory (DFS-like).
  - Good for revealing **structural equivalence**: nodes with similar roles (e.g., hubs, bridges) are visited together across communities.
  - Default:  $q = 1$  (no bias).


## Intuition: BFS vs DFS

- **BFS-like exploration** ( $q \gg 1$ , e.g.,  $q = 2$ ):
  - Walk explores outward layer by layer from starting node.
  - Stays within tight, locally-connected regions (communities).
  - Embeddings reflect community membership and local network structure.

## Intuition: BFS vs DFS

- **DFS-like exploration** ( $q \ll 1$ , e.g.,  $q = 0.5$ ):
- Walk dives deep, exploring paths and branches far from the starting point.
- More likely to reach distant or structurally similar nodes.
- Embeddings reflect structural roles: nodes acting as hubs, bridges, or peripheries cluster together even if in different communities.

# Visual Example: Transition Probabilities



## The Complete node2vec Algorithm 1/2

- **Input:** Graph  $G = (V, E, W)$ , embedding dimension  $d$ , walk length  $L$ , walks per node  $r$ , window size  $k$ , parameters  $p, q$ .
- **Step 1:** Precompute transition probabilities  $\alpha_{pq}(t, x)$  for all pairs of previous node  $t$  and candidate next node  $x$  to enable efficient sampling.
- **Step 2:** Generate walks by simulating  $r$  biased random walks of length  $L$  from each node, using the transition probabilities computed in Step 1.


## The Complete node2vec Algorithm 1/2

- **Step 3:** Train a skip-gram language model on the generated walks using word2vec, with context window  $k$  and appropriate hyperparameters (window size, negative samples, etc.).
- **Output:** Embedding matrix  $Z \in \mathbb{R}^{|V| \times d}$ , where row  $u$  is the  $d$ -dimensional embedding  $\mathbf{z}_u$ .


## Effect of $p$ and $q$ on Embeddings: Two Cliques + Bridge

- Consider a graph with two dense cliques connected by a single bridge edge.
- With **BFS-like** parameters ( $p = 1, q = 2$ ):
  - Embeddings cluster tightly by clique membership.
  - Bridge nodes stay close to their respective clique.
- With **DFS-like** parameters ( $p = 1, q = 0.5$ ):
  - Bridge nodes and other structurally similar nodes may cluster together.
  - Community boundaries are less strict in the embedding space.


## Visualizing the Effect: BFS-Like ( $p = 1, q = 2$ )



# Visualizing the Effect: BFS-Like ( $p = 1, q = 2$ )



# Visualizing the Effect: DFS-Like ( $p = 1, q = 0.5$ )



## Advantages of node2vec

- **Flexibility:** Parameters  $p$  and  $q$  allow tuning toward community detection or structural roles.
- **Efficiency:** Precomputation of transition probabilities allows fast walk generation.
- **Scalability:** Works on large networks; parallelizable walk generation.
- **Interpretability:** Resulting embeddings can be used with any ML model; clear connection to graph structure.
- **Empirical performance:** Outperforms simpler methods on node classification, link prediction, and clustering tasks.

## Practical Use: Link Prediction with node2vec Embeddings

- Compute node embeddings  $\mathbf{z}_u, \mathbf{z}_v$  for all nodes  $u, v$ .
- For each candidate pair, construct a feature vector, e.g.:

$$\mathbf{f}_{(u,v)} = [\mathbf{z}_u \circ \mathbf{z}_v, |\mathbf{z}_u - \mathbf{z}_v|, \mathbf{z}_u^\top \mathbf{z}_v]$$

- Train a classifier (e.g., logistic regression, random forest) on pairs of known edges (positive) and random non-edges (negative).
- Use learned classifier to score unknown or held-out pairs.
- This approach combines the expressive power of node2vec embeddings with classical ML.

# **DeepWalk: Learning Social Representations**

---

## What is DeepWalk?

- **DeepWalk** is a simple but powerful algorithm for learning vector embeddings for nodes in a network.
- It was introduced by Perozzi, Al-Rfou, and Skiena in 2014 as one of the first random walk-based embedding methods.
- Core idea: treat nodes in a network like words in a document, and random walks through the network like sentences in text.
- Apply a language model (skip-gram, from word2vec) to learn embeddings where nodes that frequently co-occur in walks are nearby in embedding space.

## The Key Analogy: Network as Text

- In natural language processing (NLP), word2vec learns embeddings by treating words in sentences as context for each other.
- DeepWalk applies the same principle to networks:
  - "Words" = nodes in the network
  - "Sentences" = random walks through the network
  - "Context" = nodes that appear nearby (within a window) in random walks
- Result: nodes that frequently appear together in random walks learn similar embeddings.

## Why Random Walks?

- Random walks naturally explore the local and global structure of a network.
- Starting from any node, a random walk moves to a random neighbor, then to a random neighbor of that node, and so on.
- This process discovers:
  - **Local structure:** Nodes that are close neighbors get visited together early in walks.
  - **Global structure:** Long walks eventually reach distant nodes, capturing community membership and overall graph topology.
- By sampling many random walks of various lengths, DeepWalk captures neighborhoods at multiple scales.

## Unbiased Random Walks in DeepWalk

- A random walk in DeepWalk is **unbiased**: at each step, pick a random neighbor with uniform probability.
- Formally, if current node is  $v$  with neighbors  $\Gamma(v)$ , the probability of moving to neighbor  $x$  is:

$$P(x | v) = \frac{1}{|\Gamma(v)|}$$

- This is simple and computationally fast: no need to precompute or bias transitions.
- Contrast with node2vec: node2vec introduces parameters  $p$  and  $q$  to bias walks toward BFS-like or DFS-like exploration.
- DeepWalk does not have such control; it simply explores uniformly.

## DeepWalk Algorithm: Step by Step

- **Input:** Graph  $G = (V, E)$ , embedding dimension  $d$ , walk length  $L$ , number of walks per node  $r$ , window size  $k$ .
- **Step 1:** Generate random walks.
  - For each node  $u \in V$ , simulate  $r$  unbiased random walks of length  $L$  starting from  $u$ .
  - Each walk is a sequence  $(u = v_0, v_1, \dots, v_L)$  where each  $v_{i+1}$  is a random neighbor of  $v_i$ .

## DeepWalk Algorithm: Step by Step

- **Step 2:** Collect walk data.
  - All walks together form a large corpus of sequences (like a corpus of sentences in text).
- **Step 3:** Train a skip-gram model on the walk corpus using word2vec.
- **Output:** Embedding matrix  $Z \in \mathbb{R}^{|V| \times d}$ , where row  $u$  is  $\mathbf{z}_u$ .

## Skip-Gram Model: The Learning Objective

- The skip-gram model (from word2vec) learns embeddings by maximizing the probability of context nodes given a target node.
- For a walk  $(v_0, v_1, \dots, v_L)$ , define the optimization objective:

$$\max_{\mathbf{z}_u} \sum_{u \in \text{walk}} \log P(\text{context}(u) | \mathbf{z}_u)$$

## Skip-Gram Model: The Learning Objective

- For each node  $u$  at position  $i$  in the walk, its context consists of nodes within window  $k$ :  $\{v_{i-k}, \dots, v_{i-1}, v_{i+1}, \dots, v_{i+k}\}$ .
- The conditional probability is:

$$P(v | z_u) = \frac{\exp(z_v^\top z_u)}{\sum_{w \in V} \exp(z_w^\top z_u)}$$

- In practice, negative sampling is used to avoid computing the expensive full softmax.

## Context Window: A Simple Example

- Suppose a random walk is:  $(0, 1, 2, 3, 4, 5)$  with context window  $k = 2$ .
- For node 2 at position 2 (center), the context is  $\{0, 1, 3, 4\}$  (nodes within distance 2).
- For node 2, the skip-gram model maximizes:

$$\log P(\{0, 1, 3, 4\} \mid \mathbf{z}_2)$$

- This means: given node 2's embedding  $\mathbf{z}_2$ , we want to predict that nodes 0, 1, 3, 4 are its neighbors in the walk.
- By maximizing this across all walks and all nodes, embeddings are learned so that nodes that co-occur in walks have similar (nearby) vectors.

## Negative Sampling: Efficient Training

- Computing the full softmax normalization  $\sum_{w \in V} \exp(\mathbf{z}_w^\top \mathbf{z}_u)$  is expensive for large graphs.
- **Negative sampling** approximates this by contrasting positive pairs (actual context nodes) with negative samples (random nodes).
- For each positive context node  $v$  and target  $u$ , sample a few random "negative" nodes  $w_1, \dots, w_m$  from the node distribution.

## Negative Sampling: Efficient Training

- The objective becomes:

$$\log \sigma(\mathbf{z}_u^\top \mathbf{z}_v) + \sum_{i=1}^m \mathbb{E}_{w_i} \log \sigma(-\mathbf{z}_u^\top \mathbf{z}_{w_i})$$

- Here  $\sigma$  is the sigmoid function; we maximize similarity with positive pairs and minimize with negative pairs.
- This is much faster and allows training on large networks.

## DeepWalk vs node2vec

- **Similarity:** Both use random walks to sample neighborhoods, then apply skip-gram to learn embeddings.
- **Differences:**
  - DeepWalk uses **unbiased random walks**: all neighbors equally likely.
  - node2vec uses **biased (second-order) walks** via parameters  $p$  and  $q$  to control BFS/DFS trade-off.

## DeepWalk vs node2vec

- **Flexibility:**
  - DeepWalk is simpler and faster (no precomputation of transition probabilities).
  - node2vec offers more control via  $p, q$  to tune toward community detection or structural roles.
- **Empirical performance:** Often similar; node2vec slightly better on some tasks, but the difference is often marginal.

## Why DeepWalk Works: Theoretical Connection

- DeepWalk exploits a deep connection between random walk probabilities and graph structure.
- The probability that node  $v$  appears in a random walk starting from  $u$  is related to commute times and graph partitions.
- By learning embeddings that preserve co-occurrence in random walks, DeepWalk implicitly preserves **proximity and community structure**.
- Random walks naturally concentrate on densely-connected regions (communities), so nodes in the same community are more likely to co-occur.
- This explains why DeepWalk embeddings are good for node classification and link prediction tasks.


# Hyperparameters of DeepWalk

- **Embedding dimension**  $d$ : Higher  $d$  captures more information but increases memory and training time. Typical values: 64–128.
- **Walk length**  $L$ : Longer walks explore further but are slower to compute. Typical: 40–80.
- **Number of walks per node**  $r$ : More walks give richer training signal but increase computation. Typical: 10–80.
- **Context window**  $k$ : Size of the neighborhood in skip-gram. Typical: 5–10. Larger  $k$  uses more context.
- **Negative samples**  $m$ : Number of random negatives per positive. Typical: 5–15.
- These hyperparameters must be tuned for specific


## Practical Use: Node Classification with DeepWalk

- Compute embeddings  $\mathbf{z}_u$  for all nodes using DeepWalk.
- For a downstream task like node classification:
  - Train a classifier (e.g., logistic regression, SVM, neural network) on labeled nodes using their embeddings as features.
  - Apply the trained classifier to unlabeled nodes to predict their labels.
- Similarly, for link prediction:
  - Construct pairwise features from embeddings (e.g., elementwise product, concatenation).
  - Train a binary classifier on known edges (positive) and random non-edges (negative).
  - Score unknown pairs to predict future links.

# DeepWalk on Karate Club: Example



# Comparison: Random Walks from Different Starting Nodes



## Advantages of DeepWalk

- **Simplicity:** Easy to implement and understand; no complex bias parameters.
- **Speed:** Unbiased walks are fast; no need for precomputing transition probabilities.
- **Scalability:** Random walk generation is parallelizable; works on large networks.
- **Theoretical grounding:** Connection to commute times and random walk proximity well-understood.
- **Empirical effectiveness:** Strong performance on node classification, link prediction, and clustering.
- **Flexibility:** Works on weighted, directed, and undirected graphs.

## Limitations of DeepWalk

- **No bias control:** Unbiased walks don't allow tuning toward specific structures (unlike node2vec's  $p, q$ ).
- **Hyperparameter tuning:** Multiple hyperparameters ( $d, L, r, k, m$ ) require careful tuning for best results.
- **Locality bias:** Random walks tend to stay in communities; may not capture global structure as well as other methods.
- **Computational cost:** Generating many long walks can be slow for very large graphs.
- **Stochasticity:** Results depend on random sampling; different runs may yield different (but similar) embeddings.

## DeepWalk in Context: Random Walk-Based Methods

- DeepWalk opened the door to random walk-based embedding methods.
- Extensions and improvements include:
- **node2vec**: Biased walks via  $p, q$  parameters.
- **MetaPath2Vec**: Random walks on heterogeneous networks using meta-paths.
- **Graph2Vec**: Extending to entire graphs instead of individual nodes.
- Today, graph neural networks (GNNs) often outperform these methods, but DeepWalk remains a valuable baseline and foundation.